

Fall26_SLS775. Natural Language Processing in Corpus Research

Hakyung Sung, Second Language Studies, University of Hawaii at Manoa

Last updated: September 15, 2026

Contents

1	Introduction to NLP for corpus linguistics	2
1.1	What is corpus linguistics?	3
1.2	What is NLP?	5
2	NLP in learner corpus research	9
2.1	What is learner corpus research?	9
2.1.1	Learner corpus data	9
2.1.2	Methodological approaches and research designs	10
2.1.3	Linguistic foci	11
2.1.4	Contributions to SLA	11
2.1.5	Contributions to L2 teaching and assessment	12
2.2	How can LCR be connected to NLP?	13
2.2.1	Common NLP (pre-)processes	13
2.2.2	Challenges in evaluating NLP accuracy	17
3	Corpus design and selection	20
3.1	Sampling and representativeness	20
3.1.1	An operational definition of a corpus	21
3.1.2	Size.	21
3.1.3	Principled sampling.	21
3.1.4	Representativeness.	22
3.1.5	Reusing existing corpora.	23
3.1.6	Representativeness and balance	23
3.2	Different types of corpora	24
3.3	Different types of information	24
3.4	Different ways to process corpora	26
4	Unit of analysis	29
4.1	Lexical units	29
4.1.1	Larger units in vocabulary pedagogy	30
4.1.2	Smaller units in vocabulary research	30
4.1.3	Effects on vocabulary research and pedagogy	30
4.2	Tokenization	32
4.2.1	Tokens and types	32
4.2.2	Written and spoken Language	33
4.2.3	Word boundaries across languages	33
4.2.4	Morphemes as possible units	33
4.2.5	Characters and Unicode	34
4.2.6	Subword tokenizers	34

1 Introduction to NLP for corpus linguistics

References

- [Dunn \(2022\)](#). *Natural language processing for corpus linguistics*. Cambridge University Press. Chapter 1.1: Scaling Up Corpus Linguistics. [[Section 1](#) of these notes]
- [Gries \(2026\)](#). [Corpus linguistics and psycholinguistics](#). In *International encyclopedia of language and linguistics* (3rd ed., pp. 412–416). Elsevier. [[Section 1.1](#)]
- [Meurers \(2020\)](#), Natural language processing and language learning. *Encyclopedia of applied linguistics*, 4193–4205. [[Section 1.2](#)]

Corpus linguistics and natural language processing (NLP) both investigate language through collections of naturally occurring linguistic data. However, the two fields have traditionally differed in their primary goals. Corpus linguistics generally begins with a linguistic question and uses corpus data to describe patterns of language use, while NLP develops computational methods for automatically analyzing or generating human language. NLP for corpus linguistics brings these perspectives together: computational models are used to identify, annotate, measure, and compare linguistic patterns in corpora at a scale that would be difficult to achieve through manual analysis alone.

This combination has become increasingly important as the size and diversity of available language data have expanded. Early electronic corpora, such as the Brown Corpus, contained approximately one million words, whereas contemporary corpora may contain billions of words. Such datasets create new opportunities for investigating linguistic variation, but they also make exclusively manual analysis impractical. Computational methods therefore allow corpus linguists to scale up their analyses while maintaining an explicit and, in principle, reproducible analytical procedure.

One way to organize computational problems in corpus linguistics is to distinguish between two broad types of analysis: (1) *categorization* and (2) *comparison*. Categorization involves assigning a predefined label to a linguistic unit. For example, a corpus linguist may classify a word by its part of speech, determine whether a sentence instantiates a particular grammatical construction, or identify the genre to which a document belongs. Comparison, by contrast, involves assessing the degree of similarity or relatedness between linguistic units. A researcher may examine whether two words are semantically related, whether two sentences convey similar meanings, or whether two texts share stylistic characteristics.

These two types of analysis are not rigidly tied to particular machine-learning algorithms. Note that in *supervised* learning, a model is trained on examples that have already been assigned *target labels* (in other words, annotated tags). The model learns the relationship between the input

features and the expected output and then applies that relationship to new data. In *unsupervised* learning, by contrast, the data are not provided with predefined target labels. The model instead identifies patterns, groupings, or relationships within the data.

Categorization is often implemented through supervised learning because researchers typically define the categories that a model is expected to predict. For example, a researcher may provide texts labeled as academic writing, news, or conversation and train a classifier to assign new texts to one of these registers. However, categories can also be explored through unsupervised methods. A clustering algorithm might group texts according to their linguistic similarities, after which the researcher examines whether the resulting clusters correspond to recognizable registers.

Comparison is often conducted using representations or similarity measures that do not require predefined class labels. For example, a researcher may represent two sentences as vectors and calculate their cosine similarity to estimate how closely related their meanings are. However, comparison can also involve supervised learning. A model may be trained on sentence pairs that have been labeled with degrees of semantic similarity and then used to predict the similarity of new sentence pairs.

1.1 What is corpus linguistics?

Corpus linguistics is an *empirical* approach to the study of language based on the systematic analysis of collections of linguistic data known as *corpora*. At the most basic technical level, a corpus may consist of multiple machine-readable text files accompanied by information about the texts, speakers, and contexts from which the data were collected (i.e., metadata). However, a corpus is not simply any large collection of texts. Its value as a research resource depends on how systematically the data were selected, organized, and documented. This distinction also helps clarify the difference between linguistic *data* and a *corpus*: whereas data may refer broadly to any collection of linguistic observations, a corpus is a deliberately constructed and documented dataset designed to support systematic analysis of a particular population or domain of language use.

A corpus is generally designed to represent a particular population or domain of language use, such as a group of speakers, a register, a genre, a topic, a historical period, or a language variety. Its sampling scheme should therefore reflect relevant sources of variability within the population it is intended to represent. For example, a corpus of contemporary spoken English may sample speakers from different regions, age groups, and communicative settings. Claims based on a corpus are consequently limited by the population and contexts represented in its design. We will return to this issue in our discussion of corpus *representativeness*.

Corpus linguistics is primarily concerned with the distribution of linguistic forms and functions. Researchers examine how frequently linguistic elements occur, the contexts in which they occur, the other elements with which they co-occur, and how these distributions vary across speakers, texts, registers, or time periods. These distributional patterns can be summarized through

measures (e.g., token frequency, association strength).

Corpora are commonly used in psycholinguistics as sources of quantitative measures derived from the distributional patterns of language use. These measures can serve as operationalizations of psycholinguistic constructs that are not directly observable. In other words, corpus statistics may not be able to measure cognitive representations or processes directly, but they quantify aspects of an individual's likely linguistic experience that may be relevant to how language is (learned and) used. For example, one of the most basic corpus measures is *token frequency*, or the number of times a particular linguistic form or pattern occurs. Token frequency can be related to *entrenchment*, the process through which repeated exposure strengthens a linguistic representation and allows it to become established as a cognitive routine.¹ Similarly, the frequency with which two linguistic elements occur together may contribute to the entrenchment of that combination. Corpus data also make it possible to quantify *contingency*, or the extent to which the occurrence of one linguistic element predicts the occurrence of another. Such relationships can be investigated through analyses of collocation and/or collocation, often using conditional probabilities or association measures. For example, researchers may examine whether a particular verb occurs with a given construction more frequently than would be expected from the independent frequencies of the verb and the construction. Strong contingency may facilitate the learning of form–form or form–meaning associations because the occurrence of one element provides reliable information about another.

One important advantage of corpus data is their potential for ecological validity.² Experimental studies are often conducted under highly controlled conditions using carefully selected and sometimes decontextualized stimuli. Such control is necessary for establishing causal relationships, but the resulting language may not always reflect the distributions encountered in ordinary communication. Balanced experimental designs may also expose participants to linguistic forms in proportions that differ substantially from those found in everyday language use. Corpus data complement experimental evidence by documenting how language is distributed in naturally occurring, less experimentally constrained communicative contexts.

At the same time, corpus data present several analytical challenges. Naturally occurring language is often noisy and heterogeneous, and linguistic variables are frequently highly intercorrelated. Contexts may vary considerably both within and across speakers, making it difficult to isolate the influence of any single factor. Corpus distributions are also typically highly imbalanced and approximately Zipfian: a small number of linguistic elements occur very frequently, whereas most occur only rarely. Consequently, corpus frequencies and associations must be interpreted in light of corpus composition, sampling decisions, dispersion, and the number of relevant opportunities for occurrence.

¹For example, a word, phrase, or grammatical construction that occurs frequently in language use may be recognized or processed more readily than a less frequently encountered alternative.

²Ecological validity refers to “the degree to which a given research context is indicative of, or can be generalized to, real-world environments” (Schmuckler, 2026). More details here: <https://oecs.mit.edu/pub/1y8kpgak>

Corpus linguistics is therefore not simply the automated counting of words in large text collections. It is a research approach that connects systematically sampled language data, quantitative descriptions of linguistic distributions, and theoretically motivated interpretations of language use. Its strength lies in allowing researchers to investigate how linguistic patterns emerge across large numbers of contextualized observations while remaining attentive to the limitations of the data and the procedures used to analyze them.

1.2 What is NLP?

NLP is concerned with the automated analysis and generation of human language. NLP systems process written or spoken language in order to identify linguistic properties, extract information, make predictions, or produce new language. Common NLP tasks include dividing texts into words and sentences, assigning part-of-speech (POS) categories, identifying syntactic relationships, measuring semantic similarity, translating between languages, classifying documents, answering questions, and generating text. Speech recognition and speech synthesis are also closely related to NLP, although speech processing is sometimes treated as a separate subfield.

NLP can operate at different levels of linguistic analysis. At the word level, a system may identify lemmas, morphological features, or POS tags. At the sentence level, it may analyze grammatical structure, identify errors, or represent meaning. At the text level, it may measure cohesion, classify genre, estimate proficiency, or generate a summary. NLP therefore provides methods for transforming unstructured language into representations that can be systematically searched, counted, compared, and modeled.

This role makes NLP particularly important for corpus linguistics. A corpus may contain thousands or millions of texts that cannot be analyzed manually in their entirety. NLP tools can automatically annotate these texts and extract linguistic features relevant to a research question. For example, a researcher may use an NLP pipeline to identify verbs and their subjects, calculate lexical and syntactic complexity measures for individual texts, or represent texts as vectors for classification and comparison. NLP therefore supports the scalability of corpus analysis. However, the validity of the resulting analysis depends on several factors: the capabilities and limitations of the computational method used (e.g., a simple count-based system vs. Transformer architecture), the characteristics of the corpus to which it is applied (e.g., edited standard-language texts from news sources vs. L2 learner language), and the extent to which the computational output provides a valid operationalization of the linguistic construct under investigation (e.g., text length vs. grammatical complexity).

In the context of language learning, NLP has two broad applications. The first is the analysis of *learner language*: words, sentences, and texts produced by language learners. The second is the analysis of *language for learners*: texts and other language materials that are selected, adapted, or generated to support learning.

The first broad application includes areas such as intelligent language tutoring systems, au-

omatic language assessment, and grammatical error detection. An intelligent language tutoring system may analyze a learner's response in order to provide individualized feedback, select an appropriate subsequent activity, or update a model of the learner's developing abilities. NLP is necessary in such systems because it is generally impossible to anticipate every response that a learner might produce and manually specify the appropriate feedback for each possible string. Instead, the system must abstract away from the exact wording of the response and identify more general properties of its form, meaning, and appropriateness. This requirement creates a distinctive challenge for NLP. Many general-purpose NLP systems are designed to remain robust when input contains errors/unexpected forms. A parser, for instance, is typically expected to produce a syntactic analysis even when a sentence is not fully well formed. In a language tutoring system, however, the unexpected form may itself be the information that the system needs to identify. The system must determine not only that a learner's response differs from an expected target but also what kind of difference has occurred, what it may reveal about the learner's developing language system, and what feedback would help the learner understand the problem. High correction accuracy alone is therefore not always sufficient. A system that produces an acceptable corrected sentence without identifying the linguistic source of the problem may provide limited insight into the learner's underlying difficulty and limited support for learning.

NLP is also used to annotate and analyze learner corpora with a goal to investigate broader patterns of development and variation. NLP can support error annotation, identify overuse/underuse of linguistic forms, calculate measures of complexity, accuracy, and fluency, and examine relationships between linguistic features and proficiency. It can also help researchers investigate cross-linguistic influence by identifying patterns associated with learners' first-language backgrounds. By automating at least part of this process, NLP makes it possible to analyze larger learner corpora and connect theoretical questions more directly to relevant instances in the data. This application closely reflects the role of NLP in corpus linguistics introduced in Section 1 and will be a central focus throughout this course.

The second broad application involves analyzing language that is presented to learners. For example, NLP can help identify reading materials that contain particular vocabulary or grammatical structures, estimate the difficulty or readability of texts, highlight target forms, provide links to definitions or grammatical explanations, and generate exercises or assessment items from existing materials. Such applications can give learners more individualized access to language resources and allow instructional materials to be selected according to their proficiency, goals, or current learning needs.

We briefly discussed NLP within the context of applied linguistics, but the field can be defined more broadly when viewed from the perspectives of computer science and artificial intelligence. The following definition, adapted from [Vajjala et al. \(2020, p. 6\)](#), highlights this broader perspective:

NLP is an interdisciplinary field at the intersection of computer science, artificial in-

telligence, and linguistics. It focuses on developing computational systems that can process, analyze, and understand human language. Although NLP was historically concentrated in academic institutions and research laboratories, recent advances have enabled its widespread adoption in areas such as retail, healthcare, finance, law, marketing, and human resources. Several developments have contributed to this expansion:

1. the growing availability of accessible NLP tools, techniques, and application programming interfaces (APIs);
2. improvements in generalizable and interpretable approaches to complex tasks, including open-domain dialogue and question answering;
3. increased investment in interactive consumer technologies that use language as a primary means of communication;
4. the availability of open-source datasets and standardized benchmarks; and
5. the development of datasets and language-specific models for languages beyond English and other widely digitized languages.

Chapter Summary

Corpus linguistics and NLP both use collections of language data, but they have traditionally emphasized different goals. Corpus linguistics begins with a question about language use and examines the distribution of linguistic forms and functions in (systematically constructed) corpora. NLP develops computational methods for processing, analyzing, and generating language. NLP for corpus linguistics brings these perspectives together by using computational methods to annotate, search, measure, categorize, and compare linguistic patterns across datasets that may be too large for entirely manual analysis.

Activity: Brainstorm your project for this course

You might not have a concrete research question yet. Start with a topic that you are currently studying or may be interested in studying throughout the course.

1. Briefly describe your research area or a research question that interests you. What aspect of language, language use, or language learning would you like to investigate?
2. What kind of language data could help you investigate this question? For example, would you examine learner writing, classroom interaction, interviews, social media posts, published texts, assessment responses, or another type of data?
3. Could this dataset be considered a corpus? What information would you need about its speakers, writers, texts, tasks, or contexts in order to interpret it appropriately?
4. What linguistic feature, pattern, or construct would you examine in the data? Examples might include vocabulary use, grammatical constructions, errors, syntactic complexity, semantic similarity, interactional features, or differences across learner groups.
5. Would your analysis mainly involve *categorization*, *comparison*, or both? Explain your choice using one example from your research topic.
6. What part of the analysis might NLP help automate? What would you still want a researcher to examine or verify manually?
7. Identify potential challenge(s). For example, could the corpus be unrepresentative, could the NLP tool perform poorly on the language data, or could the computational measure fail to capture the intended linguistic construct?

Group report: Select one research example from your group and briefly explain (1) the research question, (2) the proposed language data, and (3) one way that NLP could support the analysis.

2 NLP in learner corpus research

References

- Granger (2026). *Corpus-Based Research in Second Language Acquisition*. In *International encyclopedia of language and linguistics* (3rd ed.). Elsevier. [Section 2.1]
- Kyle (2021). Natural language processing for learner corpus research. *International Journal of Learner Corpus Research*, 7(1), 1-16. [Section 2.2]
- Kyle and Sung (2026). Tagging and parsing. In *International encyclopedia of language and linguistics* (3rd ed.). Elsevier. [Section 2.2]

2.1 What is learner corpus research?

Learner corpus research (LCR) is a field that uses learner corpora to investigate the characteristics and development of learner language. Learner corpus research has generally pursued two closely connected objectives. The first is *theoretical*: LCR seeks to contribute to second language acquisition (SLA) research by providing more detailed descriptions of learner language and by improving our understanding of the factors that influence it. These factors may include learners' first-language backgrounds, proficiency levels, learning contexts, task characteristics, and cognitive or affective variables. The second objective is *pedagogical*: findings from learner corpora can be used to develop teaching materials, assessment practices, and language-learning tools that more accurately address learners' needs.

2.1.1 Learner corpus data

Early learner corpora focused predominantly on L2 English. Since then, learner corpora have been developed for a much wider range of target languages. An overview of existing resources is available through the Learner Corpora around the World³ database. Nevertheless, English remains disproportionately represented, particularly in corpora of academic writing.

A basic distinction can be made between *raw* and *annotated* learner corpora. Raw learner corpora contain textual or transcribed spoken data in their original, relatively unprocessed form (usually in .txt format). They may include learner metadata (e.g., first language, age, proficiency, instructional context, or task type - usually in a separated file), but the linguistic data themselves have not been systematically enriched with additional analytical information. Annotated learner corpora include information about particular linguistic features (could be in .xlm or .conllu⁴ formats). Depending on the research purpose, annotation may cover parts of speech, morphological structure, syntactic dependencies, semantic categories, pragmatic functions, discourse features, pronunciation, disfluencies, or characteristics of spoken interaction. Some forms of annotation can be produced automatically through tools such as part-of-speech taggers, morphological analyz-

³<https://uclouvain.be/en/research-institutes/ilc/cecl/learner-corpora-around-the-world.html>

⁴<https://universaldependencies.org/format.html>

ers, and syntactic parsers. When these tools perform with sufficiently high accuracy, they make it possible to process much larger collections of learner language than would be feasible through entirely manual analysis.

Other forms of annotation continue to require substantial human judgment. Error annotation is an important example. Researchers may identify a problematic form, categorize the type of error, and sometimes reconstruct the learner's intended target form. However, both error detection and error categorization are inherently interpretive. An annotator must decide whether an unusual form is actually erroneous, what the learner may have intended, and which linguistic category best describes the problem. These judgments may be especially difficult when a sentence contains several interacting errors. For this reason, manually annotated learner corpora should be accompanied by detailed annotation manuals that define the categories and provide examples of difficult cases. Researchers should also report intra-rater or inter-rater reliability statistics when the annotation depends on human judgment. Reliability assessment does not eliminate subjectivity, but it makes the annotation process more transparent and provides evidence regarding the consistency of the resulting data.

2.1.2 Methodological approaches and research designs

The characteristics of learner corpus data have led to the development of methodological approaches specifically designed for investigating learner language. One of the best-known approaches is *Contrastive Interlanguage Analysis* (CIA; [Granger 2015](#)). CIA originally focused on comparisons between learner language and corresponding native-speaker or expert varieties. For example, researchers might compare the frequencies of particular words, grammatical constructions, collocations, or discourse markers in learner essays and expert academic writing.

Most learner corpus research has used a *cross-sectional* design, in which data are collected from different learners at one point in time. Researchers may then compare groups representing different proficiency levels or instructional stages. Cross-sectional designs are relatively practical because they do not require researchers to follow the same individuals for an extended period.

A *longitudinal* design, by contrast, collects repeated samples from the same learners over time. Longitudinal learner corpora are less common because they require sustained access to participants and must address issues such as participant attrition, changes in instructional context, and variation in elicitation tasks. Nevertheless, longitudinal learner corpus research is increasing because it provides more direct evidence of individual development.

A third approach is sometimes described as *pseudo-longitudinal*. In this design, researchers collect data from cohorts of learners at different proficiency levels at approximately the same point in time. The resulting differences do not demonstrate that particular individuals moved through the observed stages. However, they can provide detailed descriptions of different points along an assumed interlanguage continuum. This approach is popular in learner corpus research and

can inform the development of teaching materials or assessment resources designed for particular proficiency levels.

2.1.3 Linguistic foci

Learner corpus research investigates a broad range of linguistic phenomena, although some areas have received considerably more attention than others. A survey of the Learner Corpus Bibliography found that (1) lexical phenomena were examined in approximately 65% of studies, followed by (2) grammar at 46%, (3) discourse at 30%, and (4) pragmatics at 10%. Note that these categories are not mutually exclusive, so an individual study may contribute to more than one area.

Within lexical research, phraseology has been particularly prominent. Researchers have examined learners' use of collocations, lexical bundles, formulaic sequences, and associations between words. These studies may investigate whether learners use a narrower range of combinations, rely heavily on a small number of familiar expressions, or produce combinations that are grammatically acceptable but uncommon in the relevant register.

Grammatical learner corpus research includes topics such as tense and aspect, article use, modality, complementation, and syntactic construction use. Discourse and pragmatic studies frequently investigate features such as stance markers, engagement markers, connectors, discourse-organizing expressions, and strategies for expressing politeness or evaluation.

Most learner corpus studies have used written rather than spoken data. Academic writing has received especially strong attention, partly because written texts are easier to collect, store, and annotate, and partly because advanced academic writing creates important pedagogical needs. This emphasis has contributed valuable knowledge about L2 writing, but it also means that conversational language, informal registers, interactional development, and lower-proficiency production remain comparatively underrepresented.

2.1.4 Contributions to SLA

One of the original objectives of learner corpus research was to contribute to SLA theory. For a considerable period, however, the theoretical grounding of LCR remained relatively weak. Many early studies were primarily descriptive, reporting frequencies or differences without clearly connecting those observations to theories of language acquisition. This partly explains the initially limited uptake of learner corpus findings in the broader SLA community.

The situation has changed substantially (see [Tracy-Ventura and Paquot, 2021](#)). Learner corpus data now have a more established role in research on L2 acquisition, although corpus-based production data remain less central than experimental data (e.g., comprehension measures, acceptability judgments, reaction times, and online processing data). Learner corpora cannot directly reveal all of the cognitive processes underlying language acquisition. They can, however, provide large and diverse collections of naturally occurring or task-elicited learner production.

The scale and diversity of learner corpus data can make it possible to test whether findings based on a limited sample generalize across learner populations. For example, [Murakami and Alexopoulou \(2016\)](#) used large-scale learner data to challenge claims concerning a universal order of acquisition for L2 English morphemes. Their findings demonstrated that learners' first-language backgrounds substantially influenced the observed acquisition patterns. This illustrates how learner corpus evidence can qualify conclusions that may appear universal when they are based on only a small number of learner groups.

Learner corpus research has also made important contributions to the study of *linguistic complexity*. Corpus-based studies have helped researchers *operationalize* complexity by developing systematic measures of grammatical, syntactic, and lexical complexity. These measures have been used to examine how complexity develops over time and varies according to proficiency, L1, register, genre, and task complexity (e.g., [Biber et al., 2020](#); [Kyle et al., 2024](#); [Sung and Kyle, 2025](#)). A closely related area involves *predictive* studies. In addition to describing group-level differences, these studies examine which factors are associated with learners' linguistic choices or learning outcomes. For example, researchers may investigate which lexical, syntactic, phraseological, or morphological features predict human judgments of writing or speaking proficiency. Such analyses can inform theories of L2 development, but correlations between linguistic features and proficiency scores should not be interpreted as direct evidence of acquisition or causal relationships.

2.1.5 Contributions to L2 teaching and assessment

From its beginnings, learner corpus research has been strongly committed to improving L2 teaching through more learner-focused materials and pedagogical practices. Many of these resources have been designed for higher-intermediate or advanced learners, whose difficulties may involve subtle restrictions on usage, register, phraseology, and discourse organization rather than only basic grammatical errors.

One major area of application concerns *accuracy*. Learner corpus data can be used to identify recurrent grammatical errors, lexical infelicities, problematic collocations, and non-target-like phraseological patterns. These findings can then support materials intended to raise learners' awareness of common difficulties and improve grammatical, lexical, and phraseological accuracy. One example is the Louvain English for Academic Purposes Dictionary, or **LEAD** ([Granger and Paquot, 2015, 2022](#)). This web-based academic writing resource draws on analyses of academic texts produced by learners from several L1 backgrounds. It highlights cross-disciplinary academic words that learners frequently misuse, overuse, or underuse. It also presents their phraseological patterning, including common collocations and lexical bundles. The resource demonstrates how learner corpus findings can be translated into pedagogical support that goes beyond simple word definitions.

Learner corpora can also be incorporated into *data-driven learning* (DDL). In DDL activities, learners examine corpus examples and identify linguistic patterns for themselves, often with guid-

ance from the teacher. Learner-corpus-based DDL may expose students to examples of successful learner production, recurrent difficulties, or alternative revisions. Preliminary studies suggest that this form of instruction can be effective and is often received positively by students and teachers (Hartle, 2023). However, the direct use of L2 data in DDL remains relatively limited. More research is needed to determine which types of learner data are most useful, how examples should be selected and presented, and under what conditions learner-corpus-based DDL improves learning.

Learner corpus evidence may also inform proficiency-level descriptions, teaching sequences, and assessment criteria. For example, pseudo-longitudinal analyses can identify features that become more frequent, diverse, accurate, or sophisticated at different proficiency levels. These patterns may contribute to resources associated with frameworks (e.g., Common European Framework of Reference for Languages [CEFR]). However, corpus frequencies should not be converted directly into proficiency descriptors without considering task, register, learner population, and the communicative function of the feature.

2.2 How can LCR be connected to NLP?

At this point, we might have a good understanding that NLP broadly refers to the use of computational methods to analyze human language automatically (not synonymous with artificial intelligence, machine learning, or statistics, although it may draw on these approaches). NLP tools of varying levels of complexity have played an important role in the development of corpus linguistics generally and learner corpus research in particular. Relatively simple tools include concordancers, which retrieve all occurrences of a search word or expression and display them in their surrounding contexts, allowing researchers to identify and compare recurring linguistic patterns. More complex tools, such as part-of-speech (POS) taggers and syntactic parsers, automatically enrich corpus texts with grammatical and syntactic information. The following section first introduces several NLP processes commonly used to annotate and analyze learner corpora.

2.2.1 Common NLP (pre-)processes

Tokenization. Tokenization (i.e.g, the division of a text into word-like units called *tokens*) is typically the very first stage of an NLP pipeline. Although tokenization may appear straightforward in languages in which words are generally separated by spaces, even an English orthographic word may contain multiple linguistically meaningful components. Tokenization becomes more challenging in some languages (e.g., Korean), in which individual space-delimited units may regularly contain multiple morphemes. Decisions about token and morpheme boundaries can therefore become an important source of variation in corpus analysis.

Example: Tokenization in English and Korean

English sentence

She walked to school.

Word-level tokens

She |walked |to |school |. (5 tokens)

Morpheme-level units

she |walk |-ed |to |school |. (6 tokens)

(Note. Punctuation counted as a separate token)

Korean sentence

저는	학교에	갑니다.
저 + 는	학교 + 에	가 + ㅂ니다
<i>I + topic marker</i>	<i>school + locative marker</i>	<i>go + declarative ending</i>
<i>'I go to school.'</i>		

Space-delimited units

저는 |학교에 |갑니다 |. (4 tokens)

Morpheme-level units

저 |는 |학교 |에 |가 |ㅂ |니 |다 |. (7 tokens)

Lemmatization. Lemmatization reduces inflected word forms to a shared dictionary form or lemma. For example, *walks*, *walked*, and *walking* may be analyzed as forms of the lemma *walk*. Lemmatization is widely used to calculate lexical frequencies and diversity, although it is not necessary or preferable for every language or research question. Combining forms at the lemma level may obscure meaningful morphological variation, particularly when inflectional choices are themselves central to the analysis.

Example: Lemmatization and word-type counts

Example sentence

She *walks* to school. She *walked* home. She enjoys *walking*.

Word-form sequence

she |walks |to |school |she |walked |home |she |enjoys |walking

Lemmatized sequence

she |walk |to |school |she |walk |home |she |enjoy |walk

Count	Before lemmatization	After lemmatization
Tokens	10	10
Types	8	6
Word-form types:		
{ <i>she, walks, to, school, walked, home, enjoys, walking</i> } = 8		
Lemma types:		
{ <i>she, walk, to, school, home, enjoy</i> } = 6		

Tagging. Tagging commonly refers to the assignment of grammatical part-of-speech (POS) labels to words or tokens in a text. POS tagsets can be classified as either *coarse-grained* or *fine-grained*. Coarse-grained tagsets distinguish broad grammatical categories without encoding many distinctions within each category. For example, the Universal POS tagset (UPOS; Nivre et al., 2020) consists of 17 tags representing major word classes, including adjectives (ADJ), adverbs (ADV), nouns (NOUN), and verbs (VERB). Although UPOS distinguishes proper nouns from common nouns, it generally provides limited information about grammatical subclasses or inflectional forms within each word class.

Fine-grained tagsets encode more detailed, and often language-specific, distinctions within these broad POS categories. The Penn Treebank tagset (Marcinkiewicz, 1994; Santorini, 1990), for instance, contains 36 POS tags. Whereas UPOS assigns all common nouns the single tag NOUN, the Penn Treebank tagset distinguishes singular or mass nouns (NN) from plural nouns (NNS). The difference is even more apparent for verbs. UPOS generally assigns lexical verbs the single tag VERB, while the Penn Treebank tagset distinguishes six inflectional forms: base form (VB), past tense (VBD), gerund or present participle (VBG), past participle (VBN), non-third-person singular present (VBP), and third-person singular present (VBZ). The Sejong Korean tagset is a specialized tag set for tagging Korean morphemes.⁵

Figure 1 illustrates how an example sentence from the English Web Treebank (Bies et al., 2012; Silveira et al., 2014) is represented using the coarse-grained UPOS tagset and the more fine-grained Penn Treebank tagset.

	1	2	3	5	6	8	9	10	11	12	13	14	15
<i>Token</i>	His	artworks	were	selected	and	exhibited	in	the	British	Museum	in	2002	.
<i>UPOS tags</i>	PRON	NOUN	AUX	VERB	CCONJ	VERB	ADP	DET	PROPN	PROPN	ADP	NUM	PUNCT
<i>Penn tags</i>	PRP\$	NNS	VBD	VBN	CC	VBN	IN	DT	NNP	NNP	IN	CD	.

Figure 1: A comparison of UPOS tags vs. Penn Treebank tags. Adapted from Kyle and Sung (2026)

⁵See examples here: <https://nlpx12korean.github.io/UD-KSL/xpos/>

Parsing. *Parsing* refers to the automatic analysis and annotation of syntactic structure in a text. Two major approaches are commonly used: (1) *constituency parsing* and (2) *dependency parsing*. Although both approaches represent relationships among words and larger syntactic units, they differ in how sentence structure is organized and described.

(1) A *constituency parse* represents a sentence or utterance as a hierarchical tree composed of nested phrases and clauses. Words are grouped into syntactic constituents (e.g., noun phrase, verb phrases), which are in turn combined into larger structures. Figure 2 presents an example of a constituency parse tree. Although several constituency-based annotation frameworks exist, the Penn Treebank syntactic annotation scheme has been particularly influential in English-language treebanks,⁶ including the Penn Treebank (Marcinkiewicz, 1994) and OntoNotes (Weischedel et al., 2022). The Penn Treebank scheme includes labels for major phrase and clause categories, such as declarative clauses (S), noun phrases (NP), and verb phrases (VP). These labels represent the hierarchical organization of syntactic constituents within a sentence.




Figure 2: An example of constituency parse tree, visualized using *spaCy* with the *benepar_en3* parser. Adapted from Kyle and Sung (2026)

(2) A *dependency parse* represents syntactic structure as a set of directed relationships between individual words. Each relationship connects a syntactic *head* to one of its *dependents* and is assigned a label indicating the grammatical function of that relationship. Figure 3 shows an example of a dependency parse.

⁶A *treebank* is a corpus in which sentences have been annotated with syntactic structure.




Figure 3: An example of dependency parse tree with Universal Dependencies, visualized by *displaCy*. Adapted from [Kyle and Sung \(2026\)](#)

Some dependency annotation schemes are designed for particular languages, such as CLEARNLP for English ([Choi et al., 2015](#)), whereas others are intended to support consistent annotation across many languages. Universal Dependencies (UD), for example, provides a cross-linguistically applicable framework for representing grammatical relationships ([Nivre et al., 2020](#)).

Compared with constituency representations, dependency representations explicitly encode grammatical functions between words. For example, UD distinguishes several types of nominal dependents associated with predicates, including nominal subjects (*nsubj*), objects (*obj*), and oblique nominals (*obl*). It also distinguishes relationships between nouns, such as nominal modifiers (*nmod*) and appositional modifiers (*appos*). These dependency labels make it possible to identify not only which words are syntactically related, but also the grammatical function of each relationship.

2.2.2 Challenges in evaluating NLP accuracy

One of the challenges associated with these NLP procedures is that automatic annotation is usually produced through a *pipeline* of interdependent processes. Sentence segmentation may be followed by tokenization, lemmatization, tagging, and parsing, such that errors introduced at an early stage can propagate to subsequent stages and affect the final annotation. However, published accuracy estimates often evaluate individual components in isolation while assuming that all preceding steps have been completed perfectly. For example, POS-tagging accuracy may be calculated using manually corrected token boundaries rather than the automatically generated tokenization that researchers would encounter in actual applications.

NLP systems are commonly evaluated by training an algorithm on approximately 90–95% of a manually annotated corpus and testing it on the remaining 5–10%. During this evaluation, the automatic output must be aligned with the manually annotated reference data. Evaluating a single stage independently is useful for comparing systems and improving individual components, but it can create misleading expectations about performance in a complete learner corpus pipeline.

Overall accuracy may also conceal substantial variation across linguistic features. Frequent or relatively straightforward categories may approach perfect accuracy, while rare or ambiguous

categories remain poorly annotated. A single overall accuracy score therefore does not necessarily indicate whether the specific linguistic feature examined in a study has been annotated reliably.

A further issue is that NLP tools are not necessarily designed to identify the linguistic distinctions of greatest interest to SLA researchers (Meurers and Dickinson, 2017). A general-purpose parser may perform well according to its original evaluation criteria while still requiring additional processing to identify a particular construction or complexity feature. Researchers must therefore evaluate not only the general performance of a tool but also the accuracy of the specific output on which their analysis depends. This evaluation is difficult because manually annotated L2 corpora suitable for training and testing remain limited. Without appropriate reference data, researchers may have little evidence regarding how accurately a tool processes the learner population, genre, proficiency level, or linguistic feature under investigation. The followings are two illustrative studies that evaluated NLP preprocessing pipelines for L2 English (Kyle et al., 2024) and L2 Korean (Sung and Shin, 2025).

Chapter Summary

Learner corpus research uses collected samples of learner language to investigate the characteristics and development of second language use. It can contribute both to SLA research, by describing variation across learners, tasks, proficiency levels, and learning contexts, and to language teaching and assessment. Learner corpus studies may use cross-sectional, longitudinal, or pseudo-longitudinal designs and may focus on lexical, grammatical, phraseological, discourse, pragmatic, or interactional features. NLP supports learner corpus research by automating processes such as tokenization, tagging, and parsing. These processes make it possible to annotate and analyze datasets that may be too large for entirely manual analysis. However, NLP tools are typically developed and evaluated on particular languages, genres, populations, and annotation schemes. Their reported overall accuracy may therefore provide limited evidence about their performance on learner language or on the specific linguistic feature required for a study. Because errors can propagate through an NLP pipeline, researchers should evaluate the accuracy of the particular annotations, categories, or measures on which their conclusions depend.

Activity: Connecting NLP to your RQ

1. Revisit the RQ you are considering for your final project. What would you like to understand about learner language, language development, language teaching, or assessment?
2. What type of learner language data could help you investigate this question? For example, would you examine learner writing, spoken interviews, classroom interaction, assessment responses, discussion forums, or repeated samples collected over time?
3. What research design would be most appropriate: cross-sectional, longitudinal, or pseudo-longitudinal? What conclusions would this design allow you to make, and what conclusions would it not support?
4. What information about the learners, tasks, texts, or learning contexts would be necessary to interpret the corpus appropriately? Consider variables such as first-language background, proficiency, age, instructional context, genre, topic, and task conditions.
5. What linguistic feature or construct would you examine? For example, you might investigate vocabulary, phraseology, grammatical constructions, errors, syntactic complexity, discourse organization, pragmatic functions, or interactional features. How would you define and operationalize this construct?
6. What NLP process or tool could support the analysis?
7. How would you determine whether the NLP output is sufficiently accurate for your research question?
8. What part of the analysis would still require human judgment? Consider tasks such as identifying learner errors, interpreting intended meanings, resolving ambiguous annotations, or evaluating whether a computational measure represents the intended construct.

3 Corpus design and selection

References

- [Egbert et al. \(2022\)](#). Designing and evaluating language corpora: A practical framework for corpus representativeness. Cambridge University Press. [\[Section 3.1\]](#)
- [Gries \(2009\)](#). What is corpus linguistics?. In *Language and linguistics compass*. 3(5), 1225-1241. [\[Section 3.2\]](#) [\[Section 3.3\]](#) [\[Section 3.4\]](#)

Taking a broader perspective than the preceding chapters, this chapter considers how corpus linguistics has been applied across a wide range of linguistic subfields beyond learner corpus research, including general linguistics, applied linguistics, discourse analysis, sociolinguistics, and computational linguistics. Although the term *corpus* is used in broadly similar ways across these areas, introductory textbooks do not offer a single, universally shared definition. In general terms, *corpus* derives from the Latin word for “body” and refers to a body, or collection, of language data. The diversity of definitions reflects broader differences in how scholars conceptualize both the role of corpora in linguistic research and the nature of corpus linguistics as a field. Some researchers (e.g., [Teubert, 2005](#); [Stubbs, 2006](#)) consider corpus linguistics a theoretical approach to the study of language. Others define it primarily as a *methodology* for investigating linguistic research questions (e.g., [Gries, 2009](#)). From this methodological perspective, the goal of corpus linguistics is to construct and empirically analyze a sample of texts as the basis for meaningful generalizations about a language or discourse domain. A central concern is therefore whether the corpus adequately represents the domain about which the researcher intends to make claims.

3.1 Sampling and representativeness

The distinction between a *population* and a *sample* is a useful starting point for thinking about corpus design. In statistics, researchers are often interested in a large population that cannot be observed in its entirety. For example, a researcher may want to understand the reading habits of all university students, but it would usually be impossible to ask every university student directly. Instead, the researcher collects data from a smaller sample of students and uses that sample to make careful estimates about the larger population. Corpus design involves a similar logic. The population is the broader domain of language use that the researcher wants to understand (e.g., academic writing, newspaper editorials, classroom interaction, spoken conversation, etc.). The corpus is the sample: the actual collection of texts, transcripts, or language data selected for analysis. Just as a statistical sample should be selected in a principled way, a corpus should be designed so that it can support reasonable generalizations about the language-use domain of interest.

- **Population:** the broader group, domain, or type of language use that the researcher seeks to characterize.
- **Sample:** the actual collection of texts or language data selected from that population.

The population/sample distinction helps clarify why corpus design matters. If a corpus is a sample of a broader language-use domain, then the quality of the corpus depends not only on how much data it contains, but also on how that data was selected and what it is intended to represent. A very large collection of texts may still be limited if the texts were gathered only because they were easy to access. Conversely, a smaller corpus may be useful if it was carefully designed for a specific research question and target domain.

For this reason, corpus design requires researchers to ask several related questions: What broader domain of language use is the corpus supposed to represent? What kinds of texts or speakers should be included? How much data is needed? And what kinds of generalizations can reasonably be made from the resulting sample? These questions are central to how we define a corpus in the first place.

3.1.1 An operational definition of a corpus

A *corpus* can be defined as:

A large and principled sample of texts designed to represent a target domain of language use, such as a language, dialect, register, genre, or learner population.

This definition highlights three important characteristics: *size*, *principled sampling*, and *representativeness*.

3.1.2 Size.

There is no universally required minimum size for a corpus. Nevertheless, corpus size remains an important consideration because “large” is a relative concept. The amount of data required depends on several factors, including:

- the frequency and distribution of the linguistic feature being investigated;
- the nature and heterogeneity of the domain that the corpus is intended to represent; and
- the goals of the analysis.

A corpus that is sufficiently large for examining frequent lexical items may be too small for investigating rare grammatical constructions or low-frequency discourse features.

3.1.3 Principled sampling.

A corpus should not be viewed simply as a convenient collection of texts. Its design should be guided by explicit criteria concerning which texts, speakers, writers, tasks, genres, or contexts are included. These decisions should be connected to the research questions and the target domain.

3.1.4 Representativeness.

Representativeness concerns the extent to which a corpus can support better generalizations about a broader language-use domain. As discussed earlier, a corpus is not treated as the domain itself, but as a sample drawn from a larger domain of interest. For this reason, claims based on corpus data should be understood as estimates of linguistic patterns in that domain, rather than as direct facts about the domain as a whole.

A useful way to conceptualize representativeness is to distinguish among three levels: the (1) *target domain*, the (2) *operational domain*, and the (3) *corpus*. The (1) target domain refers to the full real-world domain that the researcher ultimately wants to describe. For example, a researcher may be interested in academic writing, newspaper writing, spoken classroom interaction, or argumentative essays written by second language learners. However, such domains are often too broad to observe directly in their entirety. Researchers therefore define an (2) operational domain: a more concrete and researchable version of the target domain, specified in terms of available texts, genres, registers, writers, tasks, time periods, institutions, or other relevant criteria. The (3) corpus is then the actual set of texts selected from this operational domain.




Figure 4: Three levels involved in evaluating corpus representativeness: target domain, operational domain, and corpus. Adapted from Egbert et al. (2022)

As shown in Figure 4, representativeness involves at least two related evaluations. First, researchers must ask whether the operational domain adequately captures the broader target domain. For example, if the target domain is “academic writing,” the researcher must decide which disciplines, genres, publication types, institutions, and time periods are included in the operational definition of academic writing. Second, researchers must ask whether the corpus adequately represents the operational domain. This requires attention to how texts were selected, whether important internal categories are included, and whether the sample is large and balanced enough for the intended analysis.

Representativeness should therefore be understood in relative rather than absolute terms. A corpus is not simply representative or unrepresentative in a general sense. Rather, it achieves a degree of representativeness with respect to a particular research goal, a particular domain description, and particular linguistic variables. At the same time, corpus that is appropriate for one research question may be inadequate for another. For example, a corpus designed to estimate the overall frequency of nominalizations in newspaper writing may not be adequate for comparing nominalization use across newspaper sections, historical periods, or political orientations unless

those dimensions were considered in the corpus design.

Evaluating representativeness requires attention to two broad sets of considerations:

- **Domain considerations:** whether the target domain has been clearly described, whether the operational domain is an appropriate approximation of that target domain, and whether relevant internal categories such as genres, registers, speakers, writers, tasks, contexts, institutions, or time periods are adequately represented.
- **Distributional considerations:** whether the corpus is large enough and internally varied enough to provide reliable estimates of the linguistic variables of interest. These variables may include, for example, the average frequency of nominalizations, attributive adjectives, passive constructions, lexical bundles, or particular morpheme types.

3.1.5 Reusing existing corpora.

Representativeness can be considered important when researchers reuse existing corpora. General-purpose corpora can be helpful because they can support many different studies, but their suitability cannot be assumed in advance. A corpus compiled for one purpose may not provide an adequate empirical basis for a different kind of claim. For example, a corpus may be suitable for estimating broad patterns in newspaper writing but unsuitable for comparing political editorials, sports reporting, and financial news if these internal categories were not sampled systematically. Similarly, a learner corpus may support general comparisons between L1 groups but may be less appropriate for analyzing task effects, proficiency differences, or rare grammatical constructions unless those dimensions are adequately represented.

3.1.6 Representativeness and balance

**Note.* This section adds Gries's (2009) discussion of representativeness and balance to the corpus-design framework introduced above.

Many corpora are designed with sampling issues in mind so that researchers can make reasonable generalizations from the corpus to a language as a whole, or to a specific variety, register, or discourse domain. For this reason, corpus compilers often aim to make their corpora both representative and balanced. A corpus is *representative* when it includes the major components of the language variety under investigation. For example, suppose a researcher is interested in phonological reduction in the speech of Californian adolescents. If the corpus only contains conversations among adolescents and their peers, it would not fully represent this variety of speech. This is because adolescents also speak with parents, teachers, and other adults, and these interactions may show different linguistic patterns. Ideally, these contexts should also be included. A corpus is *balanced* when the different components of the target variety are included in appropriate proportions. In other words, the amount of each type of data in the corpus should reflect its relative frequency or importance in the language variety being studied. For example, if dialogs make up approximately 65% of Californian adolescents' speech, then about 65% of the corpus should ideally consist of

dialog recordings.

Again, it is important to note, that representativeness and balance are best understood as theoretical ideals. Corpus compilers use these principles as guiding goals, but there is no single agreed-upon method for creating a corpus that is perfectly representative and perfectly balanced. In practice, corpus design always involves methodological decisions and compromises. For further discussion, Biber's work on corpus representativeness (Biber, 1990, 1993) is particularly useful, and McEnery (2006) also provide helpful discussions of corpus design in Sections A2, A8, and B1.

3.2 Different types of corpora

Corpora can differ in several important ways. These distinctions are useful because different research questions require different types of corpora. For example, a researcher interested in historical language change would need a different corpus from someone studying present-day academic writing or comparing translations across languages.

- General corpora vs. specific corpora
 - A *general* corpus aims to represent a language as a whole.
 - A *specific* corpus is intentionally restricted to a particular language use domain.
- Diachronic corpora vs. synchronic corpora
 - A *diachronic* corpus is designed to represent language change over time.
 - A *synchronic* corpus provides a snapshot of a language or variety at a particular point in time.⁷
- Monolingual corpora vs. parallel corpora
 - A *monolingual* corpus contains texts in one language or language variety.
 - A *parallel* corpus contains the same or comparable texts in multiple languages.
- Static corpora vs. dynamic corpora
 - A *static* corpus has a fixed size. Once it is compiled, no new texts are added.
 - A *dynamic* corpus, or monitor corpus, continues to grow as new materials are added over time.

These categories show that there is no single type of corpus that works for every research question. Choosing an appropriate corpus depends on what kind of language use the researcher wants to investigate: a whole language, a specific register, historical change, translation patterns, or ongoing changes in contemporary usage.

3.3 Different types of information

Corpora contain much more than isolated words. Because corpora represent language as it is used in authentic communicative situations, they include words in larger linguistic and discourse contexts, such as phrases, syntactic patterns, speaker turns, pauses, and intonation patterns in spoken

⁷This “point in time” does not necessarily mean a single year; it may cover a broader period, such as a decade.

data. In addition to the language data itself, corpus files may also include extra information that helps researchers understand, organize, and analyze the data. See Figure 5.




Figure 5: A simplified example of how metadata, markup, text, and annotation may be stored in or associated with a corpus file.

- Metadata
 - *Metadata* refers to information about the source or context of the corpus data. For example, it may include the publication date and source of a newspaper article, or information about speakers such as age, profession, education level, or number of interlocutors. Metadata may also include copyright or access information.
- Markup
 - *Markup* refers to information about the structure or physical organization of the original document. For example, markup may indicate headings, paragraph breaks, sections, page divisions, or speaker turns. Markup is generally more descriptive and objective than

linguistic annotation.

- Annotation
 - *Annotation* refers to interpretive linguistic information added to the corpus. For example, annotation may include part-of-speech tags, lemmas, syntactic dependencies, semantic categories, discourse functions, or learner error tags. Unlike metadata or markup, annotation reflects a particular linguistic analysis and may therefore be theory-dependent or debatable.

3.4 Different ways to process corpora

Once corpus files have been compiled and organized, researchers need ways to retrieve, transform, annotate, and analyze information from them. The choice of processing method depends on the format of the corpus, the complexity of the research question, and the level of control needed over the analysis. Some research questions can be addressed with ready-made corpus tools, while others require customized processing through programming.

- Corpus analysis programs
 - Many corpus analysis programs provide user-friendly graphical interfaces for common corpus-linguistic tasks.
 - These tools can generate outputs such as concordance lines, frequency lists, keyword lists, collocation tables, and n-gram lists.
 - Examples include *AntConc*, *WordSmith Tools*, and *#LancsBox*.
 - Online corpus platforms such as *Sketch Engine*⁸ provide similar functions, often with access to large preloaded corpora and additional tools such as word sketches.
- Programming languages
 - Many corpus linguists use programming languages (e.g., *Python*, *R*) to process and analyze corpus data.
 - Programming allows researchers to design customized searches, clean and restructure corpus files, extract linguistic features, calculate frequencies, normalize counts, merge metadata, visualize results, and run statistical models.
 - Programming is especially important when the analysis depends on several linked steps (e.g., filtering texts by metadata, extracting annotated features, normalizing frequencies, comparing groups, and evaluating the reliability of the resulting measures).
 - Compared with point-and-click tools, programming requires more initial training, but it provides greater flexibility, transparency, and reproducibility. Because each step can be saved as code, other researchers can inspect, repeat, modify, and extend the analysis.

In practice, these approaches can be combined. A researcher might first use a corpus analysis program to explore concordance lines, then use *Python* to clean the data, extract features,

⁸<https://www.sketchengine.eu/>

calculate normalized frequencies, and run statistical analyses. The most important point is that the processing method should match the research question, the corpus format, and the kind of evidence needed to support the study's claims.

A major advantage of programming is flexibility. Ready-made corpus tools are convenient, but they may not handle all file formats, annotation systems, or research needs. They may also make hidden decisions about how the data are processed. For example, a frequency list depends on how the program defines a “word”, but this is not always straightforward:

1. Tokenization: As we discussed earlier, this is the process of identifying word tokens in a text, which may involve many decisions. For example, researchers/programs must decide how to treat punctuation, abbreviations, hyphenated words, apostrophes, numbers, spelling variation, and URLs.
2. Headers and annotation: Corpus files may contain headers with metadata. If a program cannot distinguish the header from the main text, it may incorrectly count words in the header as part of the corpus.
3. Customized analysis: Programming allows researchers to define their own search criteria. For example, a researcher may want to retrieve all forms of a verb lemma and save each occurrence with its surrounding context. Such customized tasks may be difficult or impossible in standard concordance programs, but they can be implemented with programming.

In short, graphical tools are useful/accessible, especially for beginners and standard text files. However, programming languages give researchers greater control, flexibility, and transparency. This is especially important when dealing with research questions that require customized processing.

Chapter Summary

Corpus design/selection is a central issue in corpus linguistics because corpus-based claims depend on the relationship between the texts included in a corpus and the broader domain of language use that the researcher wants to understand. A corpus is not simply a convenient collection of texts; it is a sample designed to represent a target language use domain. For this reason, researchers need to consider corpus size, sampling procedures, and representativeness. This chapter also introduced different types of corpora, and showed that corpora contain not only words but also contextual information such as metadata, markup, and annotations, all of which can shape how corpus data are interpreted and analyzed.

Activity: Finding a corpus for your RQ

Goal: Identify a possible corpus for your research and evaluate whether it is appropriate for the claims you want to make.

1. Recall a RQ that you want to investigate using corpus data.
2. Identify the target domain of language use for your question. What broader language-use domain do you want to understand?
3. Search for an existing corpus that might be relevant to your research question. You may search corpus repositories, published studies, institutional corpus projects, or publicly available datasets.
4. Record the basic information about the corpus:
 - corpus name;
 - language or languages included;
 - size;
 - text type, genre, register, or task;
 - speakers or writers represented;
 - time period;
 - available metadata;
 - available annotation;
 - access conditions.
5. Evaluate the corpus in relation to your research question. To what extent does the corpus represent your target domain?
6. Identify possible limitations. For example, are some genres, registers, learner groups, proficiency levels, topics, time periods, or contexts missing or underrepresented?
7. Decide whether the corpus is appropriate for your research question. If it is appropriate, explain what kinds of claims it would support. If it is not appropriate, explain what kind of corpus you would need instead.
8. Share your findings with your group. Compare how different RQs require different corpus designs.

4 Unit of analysis

References

- Webb (2021). The lemma dilemma: How should words be operationalized in research and pedagogy?. In *Studies in Second Language Acquisition*. 43(5), 941–949. [Section 4.1]
- Jurafsky and Martin (2026). Words and Tokens. In *Speech and Language Processing*. [Section 4.2]

Words are often treated as the basic unit of analysis in NLP. Before counting words in a corpus, however, researchers must determine what constitutes a word. This involves two related but distinct decisions. First, they must select a theoretically appropriate *lexical unit*, such as a word type, lemma, flemma, or word family. Second, they must determine how that unit will be identified computationally in running text (i.e., *tokenization*). These decisions should not be conflated. A tokenizer segments text into computational units, but the resulting tokens do not necessarily correspond directly to the lexical units required for a particular corpus analysis.

4.1 Lexical units

Words can be defined and categorized in several ways. One important distinction concerns how different morphological forms are grouped together. Common units include *word types*, *lemmas*, *flemmas*, and *word families*, as described in Table 1.

Table 1: Lexical units

Unit	Definition	Example
Word type	Each distinct word form is treated as a separate unit.	<i>record</i> , <i>records</i> , <i>recorded</i> , and <i>unrecorded</i> are different word types.
Lemma	A headword and its inflectional forms belonging to the same part of speech.	The noun forms <i>record</i> and <i>records</i> constitute one lemma. The verb forms <i>record</i> , <i>records</i> , <i>recorded</i> , and <i>recording</i> constitute a separate lemma.
Flemma	A headword and its inflectional forms, without distinguishing between parts of speech.	The noun and verb forms <i>record</i> , <i>records</i> , <i>recorded</i> , and <i>recording</i> are grouped into one flemma.
Word family	A headword together with its inflectional and derivationally related forms.	A word family may include <i>record</i> , <i>records</i> , <i>recorded</i> , <i>recording</i> , <i>recorder</i> , <i>pre-record</i> , and <i>unrecorded</i> .

The choice of lexical unit affects estimates of vocabulary size, lexical diversity (e.g., Jarvis and Hashimoto, 2021), lexical coverage, and vocabulary knowledge. It is therefore not merely a technical preprocessing decision. It is part of the operationalization of the research construct.

4.1.1 Larger units in vocabulary pedagogy

The greatest value of larger lexical units may lie in pedagogy. Presenting headwords together with their inflectional and derivational forms may provide a shortcut to lexical development. Learning several related forms may be easier than learning the same number of morphologically unrelated words. For example, once a learner knows *record*, knowledge of English morphology may support the comprehension or learning of forms such as *recorded*, *recording*, and *recorder*. Organizing vocabulary around lemmas or word families can therefore highlight systematic relationships between words. However, knowledge of a headword does not necessarily imply complete knowledge of all its related forms. Learners may recognize a relationship between two forms without being able to understand or produce both forms accurately.

4.1.2 Smaller units in vocabulary research

Smaller units may often be more useful in research because they require fewer assumptions about learners' morphological knowledge. Treating forms separately allows researchers to distinguish knowledge of a base word from knowledge of its inflections and derivations. This does not mean that a smaller unit is always preferable. The most appropriate unit depends on the research question and the population being studied. Several factors should be considered.

Learner proficiency. For less proficient learners, smaller units may be more defensible because researchers cannot assume that knowledge of one form extends to related forms. A learner who knows *decide*, for example, may not necessarily know *decision*, *decisive*, or *indecisive*. Larger units may become more plausible for more proficient learners who have developed greater knowledge of the inflectional and derivational systems of the target language. Nevertheless, even advanced learners may not know all members of a word family.

Receptive and productive knowledge. The appropriate lexical unit may also depend on whether receptive or productive knowledge is being investigated. A learner may recognize the meaning of a derivative while reading but remain unable to produce it accurately. Similarly, being able to use one member of a word family does not ensure that other members can be used appropriately.

Research suggests that both first- and second-language learners may find it difficult to produce all of the derivatives associated with a known headword (e.g., [Nation, 2016](#); [Nation and Webb, 2011](#); [Vilkaitė-Lozdienė and Schmitt, 2020](#)). Word families may therefore be particularly problematic as units for measuring productive vocabulary knowledge.

4.1.3 Effects on vocabulary research and pedagogy

The choice of lexical unit has implications for several areas of vocabulary research and pedagogy.

Vocabulary learning. Research on L2 vocabulary learning may ask:

- To what extent does the choice of lexical unit influence estimates of vocabulary learning?
- To what extent are inflected and derived forms of known words recognized when they are encountered in meaningful contexts?
- Does knowledge of a base form facilitate the comprehension or learning of related forms?
- How do these relationships vary according to proficiency, morphological knowledge, frequency, and transparency?

One possible advantage of using larger lexical units is greater ecological validity (the same rationale we discussed earlier for using corpus data to study language as it is actually used.). In natural language use, learners encounter words in different inflected and derived forms rather than only as isolated dictionary headwords. A larger unit may therefore better represent the vocabulary demands of authentic reading or listening. The corresponding disadvantage is a lack of precision. If a learner understands *unpredictable*, a word-family analysis may treat this as knowledge of the PREDICT family. The analysis does not reveal whether the learner already knew *predict*, inferred the meaning from its morphemes, or learned *unpredictable* as a separate lexical item.

Word lists. Different word lists use different lexical units. Some lists contain word types, whereas others are organized around lemmas, flemmas, or word families. The lexical unit affects:

- the reported size of the list
- the number of forms represented by each entry
- the treatment of inflections and derivations
- estimates of lexical coverage
- comparisons between lists

A list of 2,000 word families represents considerably more surface forms than a list of 2,000 word types. These figures are therefore not directly comparable unless the unit of counting is considered.

Vocabulary assessment. One advantage of using larger lexical units in vocabulary tests is that the test can sample morphologically unrelated headwords. Testing *record*, *recorded*, and *recording* as three separate vocabulary items may partly measure knowledge of English morphology rather than knowledge of three distinct lexical meanings. Sampling one form from each word family may therefore provide an efficient estimate of the number of distinct lexical bases that a learner knows.

Lexical coverage and profiling. Lexical coverage refers to the percentage of running words in a text that are accounted for by vocabulary known to a learner. In general, the greater the proportion of known vocabulary in L2 input, the more likely the input is to be understood.

Coverage estimates depend on the lexical unit used. Consider the following forms:

develop, develops, developed, development, developmental

A word-type analysis treats these as five separate units. A lemma-based analysis groups some of the inflected forms but separates forms belonging to different parts of speech. A word-family analysis may treat all five as one unit. The same text may therefore produce different coverage estimates depending on whether the analysis uses word types, lemmas, lemmas, or word families. If a learner knows a headword but does not understand one of its derivatives, counting the entire family as known may overestimate effective lexical coverage.

4.2 Tokenization

Once researchers have selected an appropriate lexical unit, they must determine how units will be identified in corpora. This requires *tokenization*, which is generally the first stage of modern NLP processing. Tokenization refers to the process of segmenting continuous input into a sequence of units called *tokens*. For example, a *tokenizer* might represent the sentence

The students recorded the interview.

as

The | students | recorded | the | interview | .

In this representation, punctuation is treated as a separate token. The resulting tokens can subsequently be counted or used as input for POS tagging, lemmatization, parsing, and other forms of linguistic analysis.

4.2.1 Tokens and types

A *token* is an individual occurrence of a unit in a text, whereas a *type* is a distinct form among those tokens. Consider the sentence:

The student recorded the interview.

If capitalization is ignored and punctuation is counted, the sentence contains six tokens:

the | student | recorded | the | interview | .

However, it contains only five types because *the* occurs twice. The number of running tokens is conventionally represented as N , while the number of distinct types is represented as $|V|$, where V is the vocabulary of the text.

These counts depend on tokenization decisions. For example, researchers must decide whether punctuation is counted, whether a contraction such as *I'm* is treated as one token or two, and how numbers, abbreviations, and hyphenated expressions are handled. Note that, as with the operationalization of lexical units discussed in the previous section, there is no single tokenization scheme that is appropriate for every study. Tokenization decisions should instead be guided by

the language, the type of data, and the purpose of the analysis.

4.2.2 Written and spoken Language

In *written* English, spaces and punctuation provide useful indications of possible word boundaries. However, orthographic forms do not always correspond to linguistic words. For example, *I'm* may be treated as one orthographic unit or divided into the pronoun *I* and the contracted verb *'m*. *Spoken* language presents additional difficulties. Transcripts may include filled pauses such as *um* and *uh*, repetitions, false starts, word fragments, laughter, and overlapping speech. Whether these elements should be treated as words or tokens depends on the intended application. A filled pause may be excluded from a vocabulary analysis but retained in a study of fluency or interaction.

4.2.3 Word boundaries across languages

Space-based tokenization does not provide a universal definition of a word. Languages differ considerably in how word boundaries are represented in writing. For example, in *Chinese*, spaces do not normally mark word boundaries. Written forms consist of characters, and each character often represents a morpheme. However, a Chinese word may consist of more than one character, meaning that neither whitespace segmentation nor character segmentation directly identifies words. *Japanese* and *Thai* also do not consistently use spaces to mark word boundaries. At the same time, individual characters are generally too small to serve as appropriate word-level units. *Korean* uses spaces, but spaces mark boundaries between *eojeol* units rather than necessarily between individual lexical words or morphemes. An *eojeol* may contain a lexical stem together with one or more grammatical morphemes. For example, a noun may occur with a case particle, and a verb stem may occur with tense or sentence-ending morphology.

4.2.4 Morphemes as possible units

A morpheme is a meaningful subpart of a word. For example, the comparative form *longer* can be analyzed as the root *long* and the morpheme *-er*. Morphemes can be broadly divided into roots and affixes. The root carries the central lexical meaning, while affixes add grammatical or lexical information.

- *Inflectional* morphemes express grammatical information such as tense, number, case, or agreement. They are often productive, their meanings are relatively predictable, and they may be required in particular grammatical contexts.
- *Derivational* morphemes form new lexical items or change the meaning or grammatical category of a base. Their distribution and interpretation are often less predictable.
- *Clitics* behave syntactically like words but are phonologically, and sometimes orthographically, attached to another form.

Morphemes might appear to provide a useful alternative to words, but languages differ in their morphological organization. Some languages tend to express grammatical information

through separate words, while others combine several morphemes within a single word.⁹ Languages also differ in how clearly morpheme boundaries can be identified.¹⁰

4.2.5 Characters and Unicode

Characters provide a more clearly defined computational unit than words or morphemes. Before a computer can process written text, however, the characters must be represented in a form that the computer can store and interpret. Computers ultimately store information as numbers, so a shared system is needed to specify which number corresponds to each character. This system is called *character encoding*. Without a consistent encoding system, the same numerical value could be interpreted as different characters by different programs, or a character could be displayed incorrectly. This is important for multilingual text, since writing systems contain far more characters than can be represented by older encoding systems designed primarily for English.

Unicode provides a standardized inventory in which each character is assigned a unique numerical identifier, called a *code point*. UTF-8 is a widely used encoding that specifies how these Unicode code points are stored as bytes in a computer.

For most corpus research, the main practical consideration is that texts should use a consistent encoding, typically UTF-8. Inconsistent or incorrect encoding can cause characters to be displayed incorrectly or treated as different units during searching, tokenization, and frequency counting. However, character-level tokenization does not solve the linguistic problem entirely. Characters are relatively easy to identify computationally, but they are often too small to serve as useful units for representing words and meanings.

4.2.6 Subword tokenizers

Words and morphemes appear to be useful processing units because they often have relatively stable form–meaning associations. However, they are difficult to define consistently across languages. A further challenge is that the number of distinct words continues to grow as more text is observed. Consequently, a word-based NLP system will regularly encounter forms that were not included in its training data.

Modern NLP systems commonly address these problems through *subword tokenization*. A subword token may correspond to a complete word, a meaningful part of a word, or a smaller sequence of characters. Because these units can be recombined, a tokenizer can represent words that were not encountered as complete forms during training. For example, suppose a system has encountered *low*, *new*, and *newer*, but has not encountered *lower*. A strictly word-based system may treat *lower* as unknown. A subword tokenizer may instead represent it as:

low | er

⁹This distinction is commonly described as a continuum from more *analytic* languages to more *synthetic* languages.

¹⁰In more *agglutinative* systems, morphemes tend to remain relatively distinct and segmentable, whereas in more *fusional* systems, a single form may express several grammatical meanings simultaneously.

The system can therefore process the unfamiliar word using smaller units that are already part of its vocabulary.

Byte-Pair Encoding. Byte-pair encoding (BPE) is a commonly used method for learning a vocabulary of subword tokens. It begins with small units, typically individual characters or bytes, and repeatedly merges the adjacent pair that occurs most frequently in the training data. As the process continues, frequently occurring character sequences become single tokens, whereas rare words remain divided into smaller units. For example, suppose the training data contains many occurrences of words such as *low*, *lower*, and *lowest*. The tokenizer may first merge *l* and *o* into *lo*, and then merge *lo* and *w* into *low*. If *e* and *r* also occur frequently together, they may be merged into the token *er*. The word *lower* could then be represented as:

low | er

However, this segmentation is learned from frequency patterns rather than from morphological knowledge. BPE does not identify *low* as a root or *-er* as a suffix; it simply learns that these character sequences occur frequently enough to be included in its token vocabulary.

Chapter Summary

Words are often treated as basic units of corpus analysis, but what counts as a word depends on both theoretical and computational decisions. Researchers must first select an appropriate *lexical unit*. The choice of lexical unit affects estimates of downstream tasks. Researchers also determine how units will be identified in running text. *Tokenization* divides text into computational units called tokens, but there is no universal tokenization scheme. Words and morphemes provide potentially meaningful processing units, but they are difficult to define consistently across languages. Modern NLP systems therefore commonly rely on subword tokenization, which takes a more bottom-up, data-driven approach to segmenting text into processing units. In short, the central distinction is between *lexical-unit selection* and *tokenization*. Lexical-unit selection concerns which forms should be counted separately or grouped together in order to operationalize the linguistic construct being investigated. Tokenization concerns how running text is segmented into units that can be processed computationally.

Activity: Choosing a unit of analysis

1. Identify the construct that you intend to investigate.
2. Identify a possible lexical unit (e.g., word type, lemma, flemma, word family, or another linguistically motivated unit).
3. Explain why this unit is appropriate for your construct.
4. Consider how the use of a larger or smaller lexical unit might change the results. Would another unit produce a different estimate of your construct?

References

- Biber, D. (1990). Methodological issues regarding corpus-based analyses of linguistic variation. *Literary and linguistic computing*, 5(4):257–269.
- Biber, D. (1993). Representativeness in corpus design. *Literary and linguistic computing*, 8(4):243–257.
- Biber, D., Reppen, R., Staples, S., and Egbert, J. (2020). Exploring the longitudinal development of grammatical complexity in the disciplinary writing of l2-english university students. *International Journal of Learner Corpus Research*, 6(1):38–71.
- Bies, A., Mott, J., Warner, C., and Kulick, S. (2012). English web treebank. *Linguistic Data Consortium, Philadelphia, PA*.
- Choi, J. D., Tetreault, J., and Stent, A. (2015). It depends: Dependency parser comparison using a web-based evaluation tool. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 387–396.
- Dunn, J. (2022). *Natural language processing for corpus linguistics*. Cambridge University Press.
- Egbert, J., Biber, D., and Gray, B. (2022). *Designing and evaluating language corpora: A practical framework for corpus representativeness*. Cambridge University Press.
- Granger, S. (2015). Contrastive interlanguage analysis: A reappraisal. *International Journal of Learner Corpus Research*, 1(1):7–24.
- Granger, S. (2026). Corpus-based research in second language acquisition. In Nesi, H. and Milin, P., editors, *International Encyclopedia of Language and Linguistics*. Elsevier, 3 edition.
- Granger, S. and Paquot, M. (2015). Electronic lexicography goes local: Design and structures of a needs-driven online academic writing aid. *Lexicographica*, 31(1):118–141.
- Granger, S. and Paquot, M. (2022). The Louvain English for Academic Purposes Dictionary: User manual. CECL Papers 5, Centre for English Corpus Linguistics, Université catholique de Louvain, Louvain-la-Neuve.
- Gries, S. T. (2009). What is corpus linguistics? *Language and linguistics compass*, 3(5):1225–1241.
- Gries, S. T. (2026). Corpus linguistics and psycholinguistics. In Nesi, H. and Milin, P., editors, *International Encyclopedia of Language and Linguistics*, pages 412–416. Elsevier, 3 edition.
- Hartle, S. (2023). From learner corpus to data-driven learning (ddl) in eap writing. improving lexical usage in academic writing. *EuroAmerican Journal of Applied Linguistics and Languages*, 10(2):9–31.
- Jarvis, S. and Hashimoto, B. J. (2021). How operationalizations of word types affect measures of lexical diversity. *International Journal of Learner Corpus Research*, 7(1):163–194.

- Jurafsky, D. and Martin, J. H. (2026). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition with Language Models*. 3rd edition. Online manuscript released January 6, 2026.
- Kyle, K. (2021). Natural language processing for learner corpus research.
- Kyle, K. and Sung, H. (2026). Tagging and parsing. In Nesi, H. and Milin, P., editors, *International Encyclopedia of Language and Linguistics*. Elsevier, 3 edition.
- Kyle, K., Sung, H., Eguchi, M., and Zenker, F. (2024). Evaluating evidence for the reliability and validity of lexical diversity indices in l2 oral task responses. *Studies in Second Language Acquisition*, 46(1):278–299.
- Marcinkiewicz, M. A. (1994). Building a large annotated corpus of english: The penn treebank. *Using large corpora*, 273(31):5.
- McEnery, T. (2006). Corpus-based language studies: An advanced resource book. *London and New York: Routledge*.
- Meurers, D. (2020). Natural language processing and language learning. *Encyclopedia of applied linguistics*, pages 4193–4205.
- Meurers, D. and Dickinson, M. (2017). Evidence and interpretation in language learning research: Opportunities for collaboration with computational linguistics. *Language Learning*, 67(S1):66–95.
- Murakami, A. and Alexopoulou, T. (2016). L1 influence on the acquisition order of english grammatical morphemes: A learner corpus study. *Studies in Second Language Acquisition*, 38(3):365–401.
- Nation, I. S. P. (2016). *Making and Using Word Lists for Language Learning and Testing*. John Benjamins Publishing Company.
- Nation, I. S. P. and Webb, S. (2011). *Researching and Analyzing Vocabulary*. Heinle.
- Nivre, J., De Marneffe, M.-C., Ginter, F., Hajic, J., Manning, C. D., Pyysalo, S., Schuster, S., Tyers, F., and Zeman, D. (2020). Universal dependencies v2: An evergrowing multilingual treebank collection. In *Proceedings of the twelfth language resources and evaluation conference*, pages 4034–4043.
- Santorini, B. (1990). Part-of-speech tagging guidelines for the penn treebank project.
- Schmuckler, M. A. (2026). *Ecological Validity*. MIT Press. <https://oecs.mit.edu/pub/1y8kpgak>.
- Silveira, N., Dozat, T., De Marneffe, M.-C., Bowman, S. R., Connor, M., Bauer, J., and Manning, C. D. (2014). A gold standard dependency corpus for english. In *LREC*, pages 2897–2904.
- Stubbs, M. (2006). Corpus analysis: The state of the art and three types of unanswered questions. *Equinox eBooks Publishing*, pages 15–36.
- Sung, H. and Kyle, K. (2025). Usage-based analysis of l2 oral proficiency: Characteristics of argu-

- ment structure construction use. *Studies in Second Language Acquisition*, 47(4):991–1017.
- Sung, H. and Shin, G.-H. (2025). Towards robust morphosyntactic analysis of l2 korean: Evaluating and fine-tuning a korean language model. *ACM Transactions on Asian and Low-Resource Language Information Processing*, 24(11):1–21.
- Teubert, W. (2005). My version of corpus linguistics. *International journal of corpus linguistics*, 10(1):1–13.
- Tracy-Ventura, N. and Paquot, M. (2021). *The Routledge handbook of second language acquisition and corpora*. Routledge London and New York.
- Vajjala, S., Majumder, B., Gupta, A., and Surana, H. (2020). *Practical natural language processing: a comprehensive guide to building real-world NLP systems*. O’Reilly Media.
- Vilkaitė-Lozdienė, L. and Schmitt, N. (2020). Frequency as a guide for vocabulary usefulness: High-, mid-, and low-frequency words. In Webb, S., editor, *The Routledge Handbook of Vocabulary Studies*, pages 81–96. Routledge.
- Webb, S. (2021). The lemma dilemma: How should words be operationalized in research and pedagogy? *Studies in Second Language Acquisition*, 43(5):941–949.
- Weischedel, R., Pradhan, S., Ramshaw, L., Palmer, M., Xue, N., Marcus, M., Taylor, A., Greenberg, C., Hovy, E., Belvin, R., et al. (2022). *OntoNotes Release 5.0*. Borealis.