

Lab:
A survey of corpora

Hakyung Sung
University of Hawaii at Manoa

September 9, 2026

Topics to cover

- Continue working on Python
- A survey of corpora

Continue working on Assignment 1 (Deadline: 9/11)

Each exercise is worth **2 points**. Total: **10 points**.

- **Exercise 1:** Submit your .py and .json
- **Exercise 2:** Submit your .py and .json
- **Exercise 3:** Submit your .py and .json
- **Exercise 4:** Submit your .py and .json
- You may combine the .json files for Exercises 1–4 into a single file. If you do, please use a clear filename such as Exercises1–4.json so that I can easily identify it.
- **Exercise 5:** Submit *A Survey of Corpora* in any readable format (e.g., .txt, .docx, or .pdf) + Add information to the shared Google sheet
- Please use clear filenames that include the exercise number and your name, e.g., Exercise1_Hakyung_Sung.py.
- Put all files into a single .zip file and submit it to **Assignment 1** on Lamakū.

Exercise 4. Find the Longest Word in Each Line

Suppose you are given a text file called `sentences.txt`.

Example file:

```
I like corpus linguistics
Python is useful for text analysis
Natural language processing is fascinating
```

For each line, find and print the **longest word**.

Desired output:

```
linguistics
analysis
fascinating
```

► [hint!](#)

Exercise 5. Find a Corpus for Your Research

Find one corpus that could be useful for your research.

*Note that you may also download the raw corpus if you would like to continue working with it in later labs (e.g., tokenization, POS tagging).

For this lab, we will focus on **representativeness**. In approximately half a page, briefly describe:

- your research interest or possible research question;
- the corpus:
 - language(s)
 - population
 - modality / genre
 - approximate size
- why the corpus is useful for your research;
- what part of the corpus or linguistic feature(s) you might analyze.

Exercise 5. Evaluate the Corpus

Consider some of the following questions (you do not need to answer all of them):

- What population or language-use domain do you want to make claims about?
- How well does the corpus represent that population/domain?
- Are important groups, genres, tasks, proficiency levels, or time periods missing or unevenly represented?
- Is the corpus large and diverse enough for the feature(s) you want to investigate?
- What kinds of generalizations are reasonable based on this corpus?
- What claims should you avoid making?

Exercise 5. Submission

1. **Exercise 5:** Submit in any readable format (e.g., .txt, .docx, or .pdf) (1 point)
2. Add the corpus and the requested information to our **shared Google Sheet** (1 point); We will use this as a shared class repository that everyone can refer to when looking for corpora.

Appendix

Decompose the problem

Think about what needs to happen for **each line**:

1. Split the line into words.
2. Start with one word as the current longest word.
3. Look at each word one at a time.
4. Compare its length with the current longest word.
5. If it is longer, replace the current longest word.
6. Print the final result.

